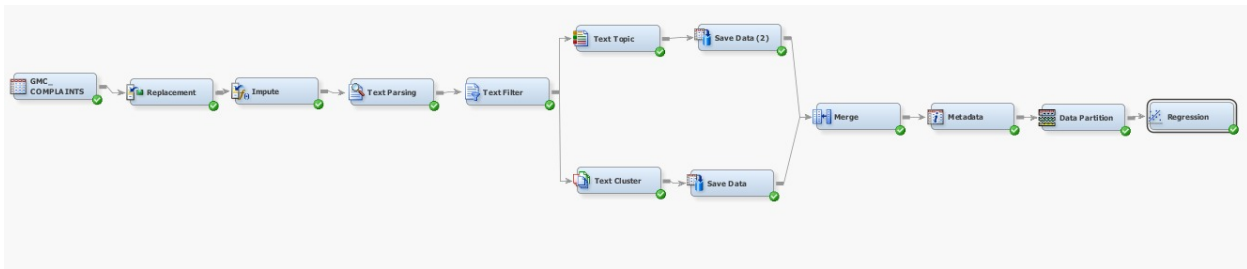


SAS Part



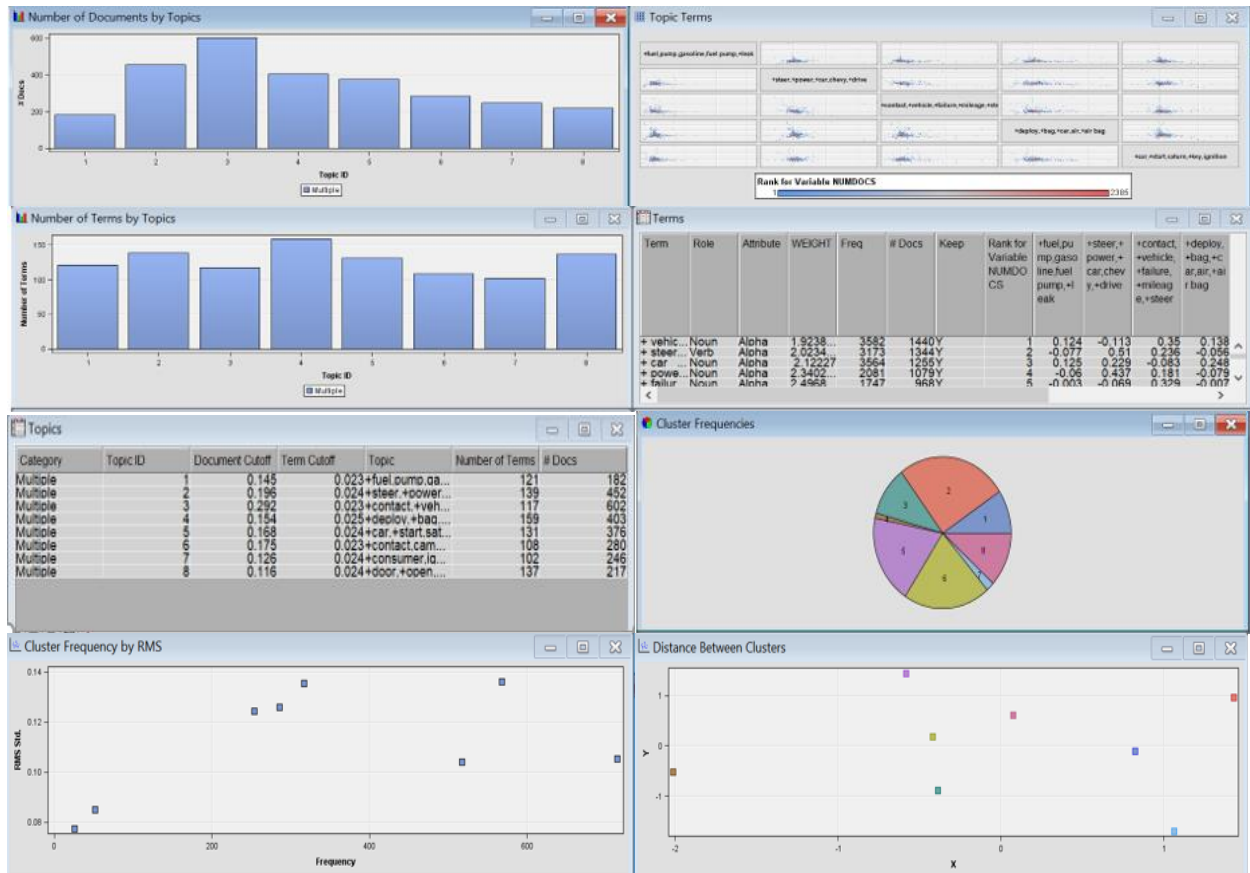
| Property | Value |
|----------------------------|-----------------------------------|
| General | |
| Node ID | TextParsing |
| Imported Data | |
| Exported Data | |
| Notes | |
| Train | |
| Variables | |
| Parse | |
| Parse Variable | Description |
| Language | English |
| Detect | |
| Different Parts of Speech | Yes |
| Noun Groups | Yes |
| Multi-word Terms | SASHELP.ENG_MULTI |
| Find Entities | None |
| Custom Entities | |
| Ignore | |
| Ignore Parts of Speech | 'Aux' 'Conj' 'Det' 'Inter' 'Part' |
| Ignore Types of Entities | |
| Ignore Types of Attributes | 'Num' 'Punct' |
| Synonyms | |
| Stem Terms | Yes |
| Synonyms | SASHELP.ENG_SYNMS |
| Filter | |
| Start List | |
| Stop List | SASHELP.ENG_STOP |
| Select Languages | |
| Report | |
| Number of Terms to Display | 20000 |

| Property | Value |
|-----------------------------|----------------------------|
| General | |
| Node ID | TextFilter |
| Imported Data | |
| Exported Data | |
| Notes | |
| Train | |
| Variables | |
| Spelling | |
| Check Spelling | No |
| Dictionary | |
| Weightings | |
| Frequency Weighting | None |
| Term Weight | Inverse Document Frequency |
| Term Filters | |
| Minimum Number of Documents | 4 |
| Maximum Number of Terms | |
| Import Synonyms | |
| Document Filters | |
| Search Expression | |
| Subset Documents | |
| Results | |
| Filter Viewer | |
| Spell-Checking Results | |
| Exported Synonyms | |
| Report | |
| Terms to View | All |
| Number of Terms to Display | 20000 |

| Property | Value |
|------------------------------|-----------|
| General | |
| Node ID | TextTopic |
| Imported Data | |
| Exported Data | |
| Notes | |
| Train | |
| Variables | |
| User Tools | |
| Term Tools | |
| Number of Single-term Topics | 0 |
| Learned Tools | |
| Number of Multi-term Topics | 8 |
| Correlated Topics | No |
| Results | |
| Topic Viewer | |

| Property | Value |
|-------------------------|---------------------|
| General | |
| Node ID | Reg |
| Imported Data | |
| Exported Data | |
| Notes | |
| Train | |
| Variables | |
| Equation | |
| Main Effects | Yes |
| Two-Factor Interactions | No |
| Polynomial Terms | No |
| Polynomial Degree | 2 |
| User Terms | No |
| Term Editor | |
| Class Effects | |
| Regression Type | Logistic Regression |
| Link Function | Logit |
| Model Options | |
| Suppress Intercept | No |
| Input Coding | Deviation |
| Model Selection | |
| Selection Model | Stepwise |
| Selection Criterion | Default |
| Use Selection Defaults | Yes |
| Selection Options | |
| Optimization Options | |
| Technique | Default |
| Default Optimization | Yes |
| Max Iterations | 0 |
| Max Function Calls | 0 |
| Maximum Time | 1 Hour |
| Convergence Criteria | |
| Uses Defaults | Yes |
| Options | |

Question 2



Question 3

| | | Predicted | |
|--------|----------|-----------|----------|
| | | Negative | Positive |
| Actual | Negative | 564 | 37 |
| | Positive | 65 | 155 |

Question 4

| | |
|-----------|------|
| Accuracy | 87.6 |
| Precision | 80.7 |
| Recall | 70.5 |
| F1 | 75.2 |

Python Part

```
import pandas as pd
import string
import nltk
import numpy as np

from AdvancedAnalytics import ReplaceImputeEncode
from AdvancedAnalytics import logreg
from sklearn.model_selection import cross_validate
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split

from nltk import pos_tag
from nltk.tokenize import word_tokenize
from nltk.stem.snowball import SnowballStemmer
from nltk.stem import WordNetLemmatizer
from nltk.corpus import wordnet as wn
from nltk.corpus import stopwords
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.decomposition import LatentDirichletAllocation

#Analyzer function
def fn_analyzer(s):
    syns = {'veh': 'vehicle', 'car': 'vehicle', 'chev': 'chevrolet', \
            'chevy': 'chevrolet', 'air bag': 'airbag', \
            'seat belt': 'seatbelt', "n't": 'not', 'to30': 'to 30', \
            'wont': 'would not', 'cant': 'can not', 'cannot': 'can not', \
            'couldnt': 'could not', 'shouldnt': 'should not', \
            'wouldnt': 'would not', 'air': 'airbag', 'bag': 'airbag'}

    s = s.lower()
    s = s.replace(',', ' . ')

#Tokenize_word
    tokens = word_tokenize(s)
    tokens = [word.replace(',', '') for word in tokens ]
    tokens = [word for word in tokens if ('*' not in word) and \
            ('''' != word) and ('``' != word) and \
            (word != 'description') and (word != 'dtype') \
            and (word != 'object') and (word != 's')]

#Mapping the synonyms
    for i in range(len(tokens)):
        if tokens[i] in syns:
            tokens[i] = syns[tokens[i]]
```

```

#stop words removal
punctuation = list(string.punctuation)+['..', '...']
pronouns = ['i', 'he', 'she', 'it', 'him', 'they', 'we', 'us', 'them']
#To add extra stop words
other =
['own','go','get','seem','say','would','regard','report','involve'\
, 'do','anoth','consumer','ve','happen','try','either','come',]
stop = stopwords.words('english') + punctuation + pronouns + other
filter_terms = [word for word in tokens if (word not in stop) and \
                (len(word)>1) and (not word.replace('.', '',1).isnumeric())\
and (not word.replace("", '',2).isnumeric())]

# Stemming
tag_words = pos_tag(filter_terms, lang='eng')
stemmer = SnowballStemmer("english")
wn_tags = {'N':wn.NOUN, 'J':wn.ADJ, 'V':wn.VERB, 'R':wn.ADV}
wnl = WordNetLemmatizer()
stemm_tokens = []
for tagged_token in tag_words :
    term = tagged_token[0]
    pos = tagged_token[1]
    pos = pos[0]
    try:
        pos = wn_tags[pos]
        stemm_tokens.append(wnl.lemmatize(term, pos=pos))
    except:
        stemm_tokens.append(stemmer.stem(term))
return stemm_tokens

#NLTK for Stop and Stem
def fn_preprocessor(s):
    #Vectorizer
    s = s.lower()
    s = s.replace(',', ' ')
    print("preprocessor")
    return(s)

def fn_tokenizer(s):
    # Tokenize
    print("Tokenizer")
    tokens = word_tokenize(s)
    tokens = [word.replace(',', '') for word in tokens ]
    tokens = [word for word in tokens if word.find('*')!=True and \
              word != "" and word !="" and word!='description' \
              and word !='dtype']
    return tokens

#column width increased
pd.set_option('max_colwidth', 32000)

```

```

df = pd.read_excel("GMC_Complaints.xlsx")

num_docs      = len(df['description'])
num_samples   = num_docs
max_features   = 1000
s_words       = 'english'
ngram         = (1,2)
max_df=0.8

discussions = []
for i in range(num_samples):
    discussions.append(("s" %df['description'].iloc[i]))

#Creating Word Frequency

cv = CountVectorizer(max_df=max_df, min_df=2, max_features=max_features,\
                    analyzer=fn_analyzer, ngram_range=ngram)
tf = cv.fit_transform(discussions)

num_topics      = 8
max_iteration    = 5
learning_offset  = 20.
learning_method  = 'online'

tf_idf = TfidfTransformer()
print("\nTF-IDF Parameters\n", tf_idf.get_params(),"\n")
tf_idf = tf_idf.fit_transform(tf)

tf_idf_vect = TfidfVectorizer(max_df=max_df, min_df=2,
max_features=max_features, analyzer=fn_analyzer, ngram_range=ngram)
tf_idf_vect = tf_idf_vect.fit_transform(discussions)
print("\nTF-IDF Vectorizer Parameters\n", tf_idf_vect, "\n")

lda = LatentDirichletAllocation(n_components=num_topics,
max_iteration=max_iteration,\
                                learning_method=learning_method, \
                                learning_offset=learning_offset, \
                                random_state=12345)

lda.fit_transform(tf_idf)
print('{:.<22s}{:>6d}'.format("Number of Reviews", tf.shape[0]))
print('{:.<22s}{:>6d}'.format("Number of Terms",      tf.shape[1]))
print("\nTopics Identified using LDA with TF_IDF")
tf_features = cv.get_feature_names()
max_words = 15
desc = []
for topic_idx, topic in enumerate(lda.components_):
    message = "Topic #d: " % topic_idx
    message += " ".join([tf_features[i]
                        for i in topic.argsort()[::-max_words - 1:-1]])

```

```

        print(message)
        print()
        desc.append([tf_features[i] for i in topic.argsort()[::-max_words -
1:-1]])

#Extracting probabilities of topic
topics = pd.DataFrame(lda.fit_transform(tf_idf))
preds = ['Year', 'make', 'model', 'crashed', 'abs', 'mileage']
df2 = pd.concat([df[preds], topics], axis=1, ignore_index=True)
df2.columns = ['Year',
'make', 'model', 'crashed', 'abs', 'mileage', '0', '1', '2', '3', '4', '5', '6', '7']

#Logistic Regression model

attribute_map = {
    'Year'      : ['I', (2003, 2011), [0, 0]],
    'make'      : ['N', ('CHEVROLET', 'PONTIAC', 'SATURN'), [0, 0]],
    'model'     : ['N', ('COBALT', 'G5', 'HHR', 'ION', 'SKY', 'SOLSTICE'), [0, 0]],
    'crashed'   : ['B', ('N', 'Y'), [0, 0]],
    'abs'       : ['B', ('N', 'Y'), [0, 0]],
    'mileage'   : ['I', (0, 200000), [0, 0]],
    '0'         : ['I', (0, 1), [0, 0]],
    '1'         : ['I', (0, 1), [0, 0]],
    '2'         : ['I', (0, 1), [0, 0]],
    '3'         : ['I', (0, 1), [0, 0]],
    '4'         : ['I', (0, 1), [0, 0]],
    '5'         : ['I', (0, 1), [0, 0]],
    '6'         : ['I', (0, 1), [0, 0]],
    '7'         : ['I', (0, 1), [0, 0]],
}
varlist = ['crashed']
rie = ReplaceImputeEncode(data_map=attribute_map, \
                           nominal_encoding='one-hot',
                           interval_scale = None, drop=False, display=False)
encoded_df = rie.fit_transform(df2)
X = encoded_df.drop(varlist, axis=1)
y = encoded_df[varlist]
np_y=np.ravel(y)

max_f1 = 0
List=[.1,1,10,100]
score_list = ['accuracy', 'recall', 'precision', 'f1']
for c in List:
    print("\nRegularization Parameter: ", c)
    lgr = LogisticRegression(C=c, tol=1e-8, max_iteration=1000)
    lgr.fit(X, np_y)
    scores = cross_validate(lgr, X, np_y, \
                             scoring=score_list, return_train_score=False, \
                             cv=10)
    print("{:.<13s}{: >6s}{: >13s}".format("Metric", "Mean", "Std. Dev."))

```

```

    for s in score_list:
        var = "test_"+s
        mean = scores[var].mean()
        std = scores[var].std()
        print("{:.<13s}{:>7.4f}{:>10.4f}".format(s, mean, std))
        if mean > max_f1:
            max_f1 = mean
            best_predictor = c
print("\nBest based on F1-Score")
print("Best Regularization Parameter = ", best_predictor)

X_train, X_valid, y_train, y_valid= \
train_test_split(X,y,test_size = 0.3, random_state=7)

np_y_train = np.ravel(y_train)
np_y_valid = np.ravel(y_valid)

lr = LogisticRegression(C=best_predictor, tol=1e-8, max_iteration=1000)
lr.fit(X_train,np_y_train)

logreg.display_coef(reg,21,2,X_train.columns)

logreg.display_binary_split_metrics(reg,X_train,np_y_train,X_valid,np_y_valid
)

```

Outputs

Metrics

| | | |
|-----------------------------|--------|--------|
| Accuracy..... | 0.8191 | 0.7881 |
| Precision..... | 0.8818 | 0.8554 |
| Recall (Sensitivity)... | 0.3573 | 0.3047 |
| F1-score..... | 0.5085 | 0.4494 |
| MISC (Misclassification)... | 18.1% | 21.2% |
| class 0..... | 1.7% | 2.0% |
| class 1..... | 64.3% | 69.5% |

Confusion Matrix

Validation

| Confusion Matrix | Class 0 | Class 1 |
|------------------|---------|---------|
| Class 0..... | 576 | 12 |
| Class 1..... | 162 | 71 |